BatonIndex

Leadership Succession Intelligence

A BATONINDEX METHODS BRIEF FOR SEARCH PROFESSIONALS

A companion deep-dive to The Graduate College Dean Pipeline (July 2026)

The Graduate Deanship Clock

How long the seat really lasts, when exit risk peaks, and the text-mined hiring signature behind who a public research university promotes into it.

277	7 yrs	42%
graduate deans tracked	hazard-model median tenure	promoted from a feeder role

Basis: 277 graduate-dean appointments at US public research universities identified from official leadership pages (113 completed, 111 sitting, both with a confirmed start year). 305 associate/assistant/interim bench rows tracked separately. August 2026.

Executive summary

- The naive number is wrong, and we can show why: across 113 completed tenures the median looks like 4 years. But that sample systematically excludes every long-serving dean who is still sitting — they haven't generated a "completed" row yet. Correcting for that with a Kaplan–Meier estimator puts the true median tenure at 7 years.
- Exit risk is not flat — it has a window: the year-by-year hazard rate climbs through the first six years, then holds at an elevated 17–29% band in years 7–10 before the sample thins out. That band is the practical watch-window for succession planning, not the mean or the median alone.
- This is an internal-ladder office: 42% of deans were promoted from an explicit associate/vice-provost/interim feeder title (25% from "associate dean" specifically). Arizona State's own succession chain is the archetype (page 4).

How long the job looks like it lasts

Take only the finished tenures on record — 113 of them, now with national coverage instead of one region — and the graduate deanship looks short: a median of 4 years, a mean of 5.4, and a long right tail out to 24 years. This is the number a simple tenure-length query returns, and it is the number most trackers stop at.

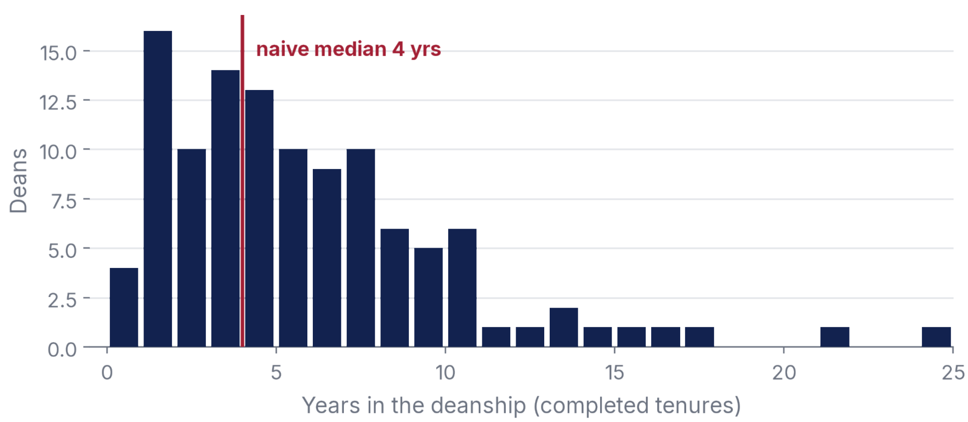

Figure 1. Distribution of 113 completed tenures nationwide. Two records with an unrecorded start year were corrected out rather than estimated (see Method, page 5).

Why a histogram alone misreads this office

A completed-tenure histogram has a built-in bias: at any snapshot in time, the deans who have served the longest are disproportionately likely to still be sitting — they simply haven't left yet, so they never enter the "completed" sample at all. The fix is standard in survival analysis and rare in leadership-search research: treat every sitting dean as a right-censored observation (still "at risk" of leaving as of their current tenure-to-date) alongside the 113 completed spells, then build a discrete-time life table across all 224 deans with a known start date.

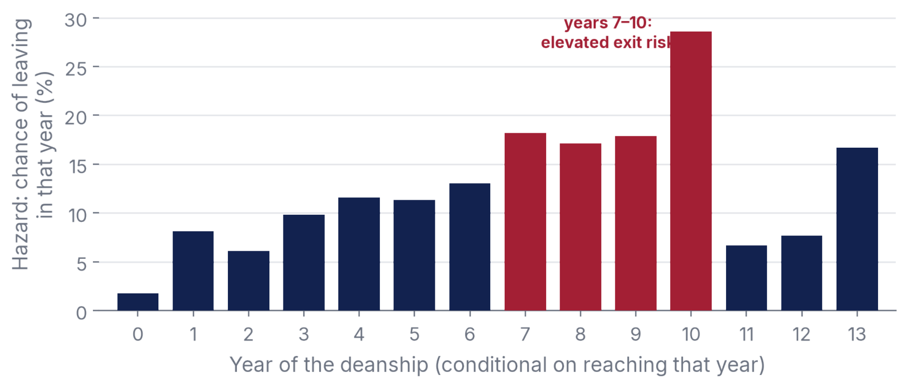

Figure 2. Conditional hazard h(t) = departures in year t ÷ deans still in the seat entering year t, computed on the full life table (completed + censored sitting deans). Years beyond 13 are omitted — the at-risk set drops under 12 and the hazard becomes too noisy to read.

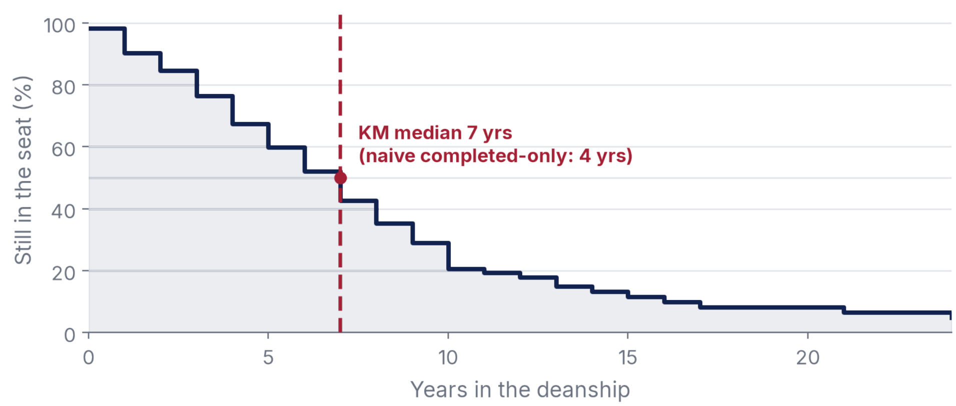

Figure 3. Kaplan–Meier survival curve, S(t) = the share of deans still in the seat after t years. The median crossing (where 50% have exited) lands at year 7, a full 3 years past the naive completed-only estimate on the previous page.

Who gets the job

We mined free-text prior-title and career-background fields across all 277 deans (not just the ones with a usable start date) because that is where this office's hiring signature actually lives — the dataset's structured discipline flags are mostly unpopulated for this index. 42% were promoted from a title that already sits one rung below the seat: associate or assistant dean, vice or associate provost, or an interim appointment. Only 20% arrived as a lateral hire from another deanship, and just 9% jumped straight from a department chair.

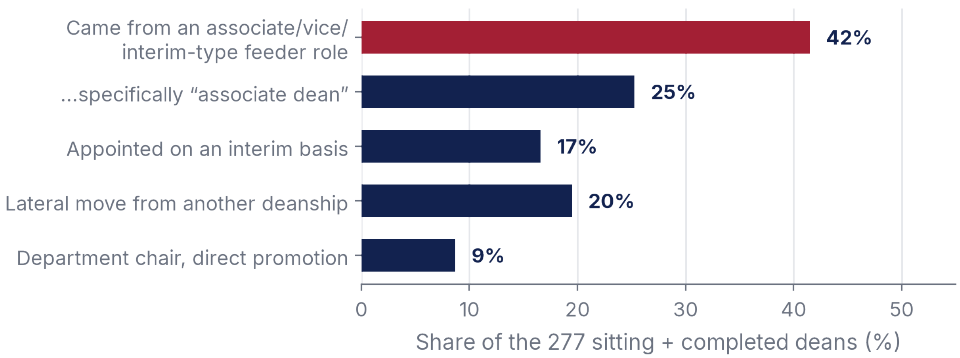

Figure 4. Prior-title keyword matches across all 277 deans (categories overlap; a title can match more than one pattern).

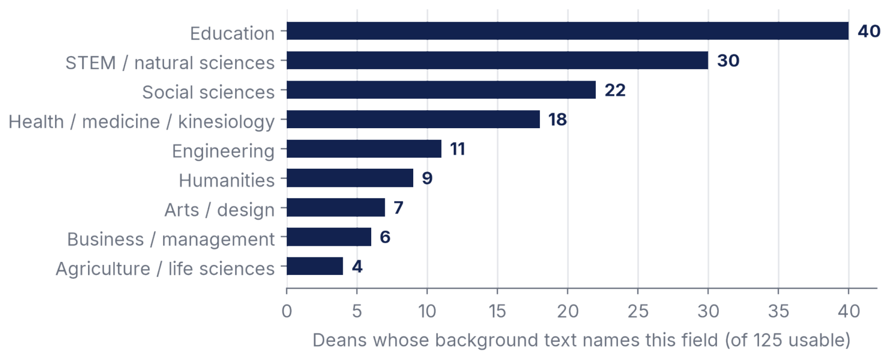

Figure 5. Disciplinary field named in free-text career-background notes (125 of 277 deans have usable text). Education and STEM lead, but no single field clears one-third of the tagged sample — the seat is genuinely discipline-agnostic.

Case study: Arizona State University's own succession chain

Every one of ASU's last four graduate-dean transitions was an internal promotion from an adjacent office — Associate VP for Graduate Support Programs (2013), Dean of Social Sciences (2020), Executive Vice Provost of the Academic Enterprise (2026, interim) — not an outside search. It is the feeder-role pattern in Figure 4, run at a single institution: grow the successor in the office next door.

What this means for a search

- Plan the reopening on the hazard curve, not the median alone: a dean entering year 7 is entering the highest-risk band for departure through year 10 — treat that window as a live pipeline stage, not a settled seat.
- Recruit from the office next door: associate/assistant deans, vice and associate provosts, and sitting interims account for four in ten placements. Build the bench list before the search opens.
- Do not screen on discipline: the widest single field (education) still accounts for well under a third of tagged appointments. Screen for cross-college administrative range instead.
- Use survival math, not a spreadsheet median, when you brief a client: “median tenure” computed on completed rows alone understates how long a strong incumbent actually stays — by several years, per Figure 3.

Method and limits

This brief extends the July 2026 Graduate College Dean Pipeline brief’s dataset (277 top-level graduate deans, now with national rather than Western-only date coverage) with two additions: a discrete-time hazard/Kaplan–Meier survival model, and a keyword-based text-mining pass over free-text prior-title and career-background fields (the index’s structured discipline and career flags are largely unpopulated). 53 records were excluded from the tenure and hazard analysis for missing a start date, including two with a recorded start year of zero — a data artifact corrected by exclusion rather than estimation, following the same rule the July brief used. The hazard curve is truncated at year 13, where the at-risk set drops under 12 deans and single departures start swinging the rate implausibly. Feeder-role and discipline shares (Figures 4–5) are pattern-matched from free text, not a hand-coded taxonomy, and are best read as directional signal, not a census. Percentages describe this sample, not a claim about every US graduate school.